


ADDISON WESLEY DATA & ANALYTICS SERIES



DEEP LEARNING ILLUSTRATED

A Visual, Interactive Guide to Artificial Intelligence



JON KROHN

with **GRANT BEYLEVELD** and **AGLAÉ BASSENS**

About This E-Book

EPUB is an open, industry-standard format for e-books. However, support for EPUB and its many features varies across reading devices and applications. Use your device or app settings to customize the presentation to your liking. Settings that you can customize often include font, font size, single or double column, landscape or portrait mode, and figures that you can click or tap to enlarge. For additional information about the settings and features on your reading device or app, visit the device manufacturer's Web site.

Many titles include programming code or configuration examples. To optimize the presentation of these elements, view the e-book in single-column, landscape mode and adjust the font size to the smallest setting. In addition to presenting code and configurations in the reflowable text format, we have included images of the code that mimic the presentation found in the print book; therefore, where the reflowable format may compromise the presentation of the code listing, you will see a "Click here to view code image" link. Click the link to view the print-fidelity code image. To return to the previous page viewed, click the Back button on your device or app.

Praise for *Deep Learning Illustrated*

"Over the next few decades, artificial intelligence is poised to dramatically change almost every aspect of our lives, in large part due to today's breakthroughs in deep learning. The authors' clear visual style provides a comprehensive look at what's currently possible with artificial neural networks as well as a glimpse of the magic that's to come."

—Tim Urban, writer and illustrator of *Wait But Why*

"This book is an approachable, practical, and broad introduction to deep learning, and the most beautifully illustrated machine learning book on the market."

—Dr. Michael Osborne, Dyson Associate Professor in Machine Learning, University of Oxford

"This book should be the first stop for deep learning beginners, as it contains lots of concrete, easy-to-follow examples with corresponding tutorial videos and code notebooks. Strongly recommended."

—Dr. Chong Li, cofounder, Nakamoto & Turing Labs; adjunct professor, Columbia University

"It's hard to imagine developing new products today without thinking about enriching them with capabilities using machine learning. Deep learning in particular has many practical applications, and this book's intelligible clear and visual approach is helpful to anyone who would like to understand what deep learning is and how it could impact your business and life for years to come."

—Helen Altshuler, engineering leader, Google

"This book leverages beautiful illustrations and amusing analogies to make the theory behind deep learning uniquely accessible. Its straightforward example code and best-practice tips empower readers to immediately apply the transformative technique to their particular niche of interest."

—Dr. Rasmus Rothe, founder, Merantix

"This is an invaluable resource for anyone looking to understand what deep learning is and why it powers almost every automated application today, from chatbots and voice recognition tools to self-

driving cars. The illustrations and biological explanations help bring to life a complex topic and make it easier to grasp fundamental concepts.”


—*Joshua March, CEO and cofounder, Conversocial; author of Message Me*

“Deep learning is regularly redefining the state of the art across machine vision, natural language, and sequential decision-making tasks. If you too would like to pass data through deep neural networks in order to build high-performance models, then this book—with its innovative, highly visual approach—is the ideal place to begin.”

—*Dr. Alex Flint, roboticist and entrepreneur*



Deep Learning Illustrated



Visit informit.com/awdataseries for a complete list of available publications.

The **Pearson Addison-Wesley Data & Analytics Series** provides readers with practical knowledge for solving problems and answering questions with data. Titles in this series primarily focus on three areas:

1. **Infrastructure:** how to store, move, and manage data
2. **Algorithms:** how to mine intelligence or make predictions based on data
3. **Visualizations:** how to represent data and insights in a meaningful and compelling way

The series aims to tie all three of these areas together to help the reader build end-to-end systems for fighting spam; making recommendations; building personalization; detecting trends, patterns, or problems; and gaining insight from the data exhaust of systems and user interactions.



Make sure to connect with us!
informit.com/socialconnect



Deep Learning Illustrated

A Visual, Interactive Guide to Artificial Intelligence

Jon Krohn
with Grant Beyleveld
and Aglaé Bassens

◆◆ Addison-Wesley

Boston • Columbus • New York • San Francisco • Amsterdam • Cape Town
Dubai • London • Madrid • Milan • Munich • Paris • Montreal • Toronto • Delhi • Mexico City
São Paulo • Sydney • Hong Kong • Seoul • Singapore • Taipei • Tokyo

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and the publisher was aware of a trademark claim, the designations have been printed with initial capital letters or in all capitals.

The authors and publisher have taken care in the preparation of this book, but make no expressed or implied warranty of any kind and assume no responsibility for errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

For information about buying this title in bulk quantities, or for special sales opportunities (which may include electronic versions; custom cover designs; and content particular to your business, training goals, marketing focus, or branding interests), please contact our corporate sales department at corpsales@pearsoned.com or (800) 382-3419.

For government sales inquiries, please contact governmentsales@pearsoned.com.

For questions about sales outside the U.S., please contact intlcs@pearson.com.

Visit us on the Web: informit.com/aw

Library of Congress Number: 2019938790

Copyright © 2020 Pearson Education, Inc.

Cover art by Aglaé Bassens

All rights reserved. This publication is protected by copyright, and permission must be obtained from the publisher prior to any prohibited reproduction, storage in a retrieval system, or transmission in any form or by any means, electronic, mechanical, photocopying, recording, or likewise. For information regarding permissions, request forms and the appropriate contacts within the Pearson Education Global Rights & Permissions Department, please visit www.pearson.com/permissions/.

ISBN-13: 978-0-13-511669-2

ISBN-10: 0-13-511669-4

3 20

◆
For Gigi
—Jon
◆

Contents

Figures

Tables

Examples

Foreword

Preface

Acknowledgments

About the Authors

I Introducing Deep Learning

1 Biological and Machine Vision

Biological Vision

Machine Vision

The Neocognitron

LeNet-5

The Traditional Machine Learning Approach

ImageNet and the ILSVRC

AlexNet

TensorFlow Playground

Quick, Draw!

Evaluation Only. Created with Aspose.Words. Copyright 2003-2026 Aspose Pty Ltd.

Summary

2 Human and Machine Language

Deep Learning for Natural Language Processing

Deep Learning Networks Learn Representations Automatically

Natural Language Processing

A Brief History of Deep Learning for NLP

Computational Representations of Language

One-Hot Representations of Words

Word Vectors

Word-Vector Arithmetic

word2viz

Localist Versus Distributed Representations

Elements of Natural Human Language

Google Duplex

Summary

3 Machine Art

A Boozy All-Nighter

Arithmetic on Fake Human Faces

Style Transfer: Converting Photos into Monet (and Vice Versa)

Make Your Own Sketches Photorealistic

Creating Photorealistic Images from Text

Image Processing Using Deep Learning

Summary

4 Game-Playing Machines

Deep Learning, AI, and Other Beasts

Artificial Intelligence

Machine Learning

Representation Learning

Artificial Neural Networks

Deep Learning

Machine Vision

Natural Language Processing

Three Categories of Machine Learning Problems

Supervised Learning

Unsupervised Learning

Reinforcement Learning

Deep Reinforcement Learning

Video Games

Board Games

AlphaGo

- AlphaGo Zero
- AlphaZero
- Manipulation of Objects
- Popular Deep Reinforcement Learning Environments
 - OpenAI Gym
 - DeepMind Lab
 - Unity ML-Agents
- Three Categories of AI
 - Artificial Narrow Intelligence
 - Artificial General Intelligence
 - Artificial Super Intelligence
- Summary

II Essential Theory Illustrated

5 The (Code) Cart Ahead of the (Theory) Horse

- Prerequisites
- Installation
- A Shallow Network in Keras
 - The MNIST Handwritten Digits
 - A Schematic Diagram of the Network
 - Loading the Data
 - Reformatting the Data
 - Designing a Neural Network Architecture
 - Training a Neural Network Model
- Summary

6 Artificial Neurons Detecting Hot Dogs

- Biological Neuroanatomy 101
- The Perceptron
 - The Hot Dog / Not Hot Dog Detector
 - The Most Important Equation in This Book
- Modern Neurons and Activation Functions
 - The Sigmoid Neuron
 - The Tanh Neuron
 - ReLU: Rectified Linear Units
- Choosing a Neuron
- Summary
- Key Concepts

7 Artificial Neural Networks

- The Input Layer
- Dense Layers

- A Hot Dog-Detecting Dense Network
 - Forward Propagation Through the First Hidden Layer
 - Forward Propagation Through Subsequent Layers
- The Softmax Layer of a Fast Food-Classifying Network
- Revisiting Our Shallow Network
- Summary
- Key Concepts

8 Training Deep Networks

- Cost Functions
 - Quadratic Cost
 - Saturated Neurons
 - Cross-Entropy Cost
- Optimization: Learning to Minimize Cost
 - Gradient Descent
 - Learning Rate
 - Batch Size and Stochastic Gradient Descent
 - Escaping the Local Minimum
- Backpropagation
- Tuning Hidden-Layer Count and Neuron Count
- An Intermediate Net in Keras
- Summary
- Key Concepts

9 Improving Deep Networks

- Weight Initialization
 - Xavier Glorot Distributions
- Unstable Gradients
 - Vanishing Gradients
 - Exploding Gradients
 - Batch Normalization
- Model Generalization (Avoiding Overfitting)
 - L1 and L2 Regularization
 - Dropout
 - Data Augmentation
- Fancy Optimizers
 - Momentum
 - Nesterov Momentum
 - AdaGrad
 - AdaDelta and RMSProp
 - Adam
- A Deep Neural Network in Keras
- Regression

TensorBoard
Summary
Key Concepts

III Interactive Applications of Deep Learning

10 Machine Vision

Convolutional Neural Networks
The Two-Dimensional Structure of Visual Imagery
Computational Complexity
Convolutional Layers
Multiple Filters
A Convolutional Example
Convolutional Filter Hyperparameters
Pooling Layers
LeNet-5 in Keras
AlexNet and VGGNet in Keras
Residual Networks

This document was truncated here because it was created in the Evaluation Mode.



ASPOSE
Your File Format APIs